

# Data Mining Algorithms

## Parallelizing K-Means with MPI

Eleonora Renz, 239020

# 1. Introduction

# Objectives

- Parallelize the most popular [1] clustering algorithm using MPI
- Reduce computational time
- Gain deeper understanding of MPI, HPC, and C

[1] P. Berkhin. 2006. A Survey of Clustering Data Mining Techniques. Springer Berlin Heidelberg, Berlin, Heidelberg, 25–71.

# State of the Art

- Well studied algorithm
- Parallelization through MPI & OpenMP  
→ No code, nor data found to compare this solution to

## 2. Problem Analysis

K-Means

# Pseudocode

---

**Algorithm 1** Serial k-means

---

- 1: Set initial centroids  $k = \langle k_1, k_2, \dots, k_k \rangle$
  - 2: Calculate distance of each point to the  $k$  centroids
  - 3: Assign each point to a cluster based on the shortest distance
  - 4: Calculate new centroids based on cluster members
  - 5: If  $k$  and  $k'$  are different
  - 6: Set  $k$  to be  $k'$  and repeat from step 2
  - 7: Else stop
-

# 3. Parallel Solution Design

K-Means

# Pseudocode

---

**Algorithm 3** Parallel k-means

---

- 1: Set initial centroids  $k = \langle k_1, k_2, \dots, k_k \rangle$
  - 2: Split data amongst  $s$  processes into  $s$  subgroups of equal size
  - 3: **for** every subgroup **do**
  - 4:     Calculate distance of each point to the  $k$  centroids
  - 5:     Assign each point to a cluster based on the shortest distance
  - 6: **end for**
  - 7: Gather all points with their cluster associations
  - 8: Calculate new centroids based on cluster members
  - 9: If  $k$  and  $k'$  are different
  - 10: Set  $k$  to be  $k'$  and repeat from step 3
  - 11: Else stop and return all points with their cluster association
-



# Parallelize with MPI

In the root process:

- Read data
- Set k initial centroids

Use collective communication calls to:

- Scatter data points across different processes with MPI\_Scatter  
→ custom MPI\_Datatype was defined for sending struct data
- Broadcast initial centroids to all processes

In each process:

- Find the euclidean distance between all data points in the subset with the k centroids
- Assign points to a cluster based on shortest distance to a centroid

Use collective communication call to:

- Gather data points with their new cluster association in the root process with MPI\_Gather

```
typedef struct point
{
    double x[6]; // data points
    int cluster; // cluster association
} point;
```

# Parallelize with MPI

In the root process:

- Re-calculate centroids based on cluster members of gathered data
- Check if centroids have moved

If centroids have moved, use collective communication calls to:

- Broadcast new centroids with MPI\_Bcast
- Repetition of previous points for all processes up until a max iteration limit

Else, if centroids have not moved, use collective communication calls to:

- Break the loop in each process by broadcasting an indicating variable with MPI\_Bcast

# 4. Benchmarking

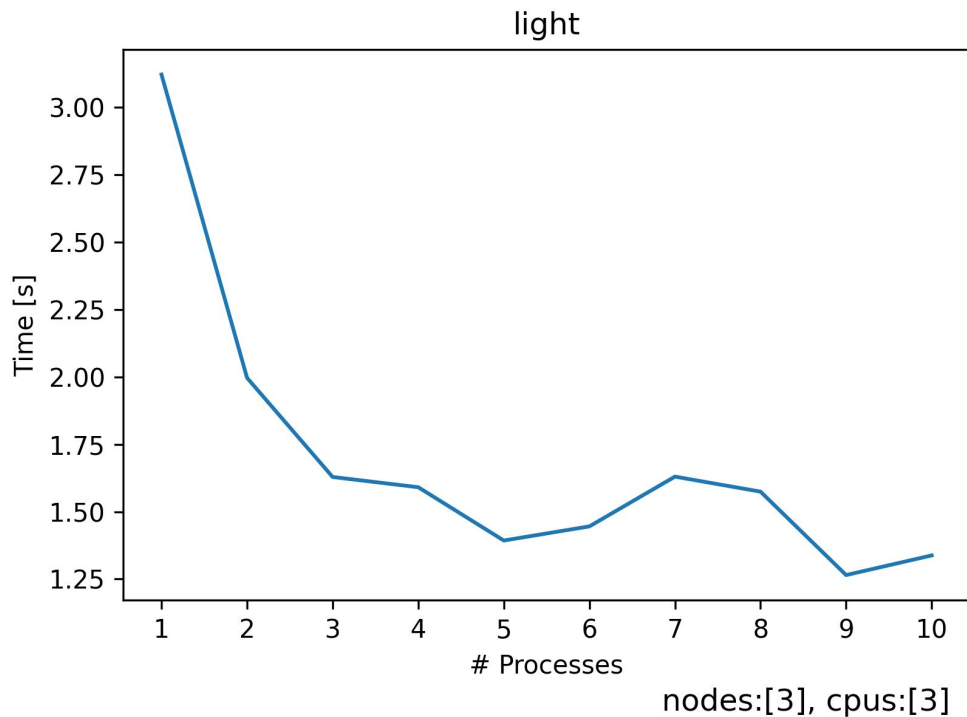
K-Means

# Benchmark Experiments


1. Light data set size (8950 rows, 6 columns) with a fixed size of nodes and CPUs while increasing the number of processes used.
  2. Heavy data set size (17950 rows, 6 columns) with a fixed size of nodes and CPUs while increasing the number of processes used.
  3. Heavy data set size with increasing nodes along number of processes while keeping CPUs fixed.
  4. Heavy data set size with increasing CPUs along number of processes while keeping nodes fixed.
- No benchmarking on I/O
  - Do 10 runs of every job and take median result

| Experiment  | Dataset size | Nodes | CPUs | Processes |
|-------------|--------------|-------|------|-----------|
| Light       | 501200B      | 3     | 3    | 1-30      |
| Heavy       | 1005200B     | 3     | 3    | 1-30      |
| Heavy Nodes | 1005200B     | 1-30  | 1    | 1-30      |
| Heavy CPUs  | 1005200B     | 1     | 1-30 | 1-30      |

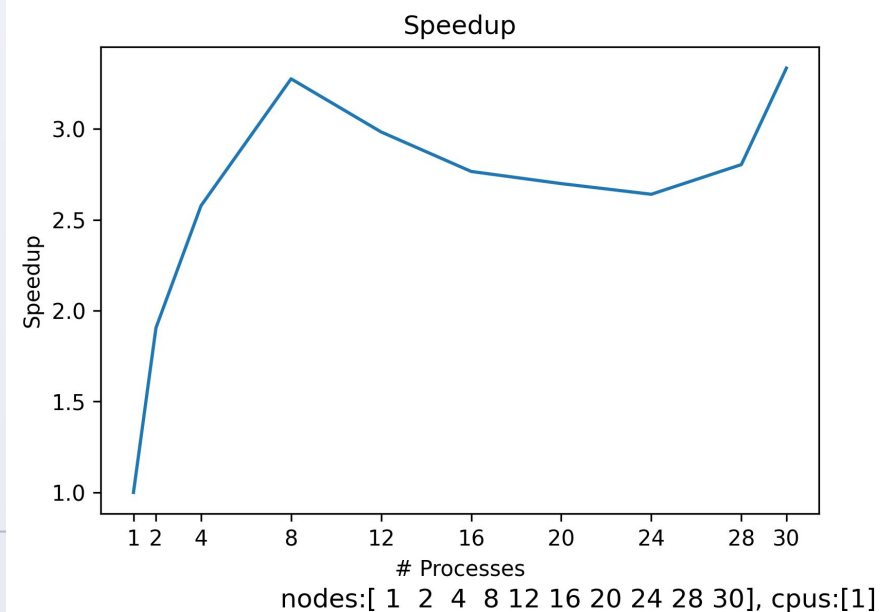
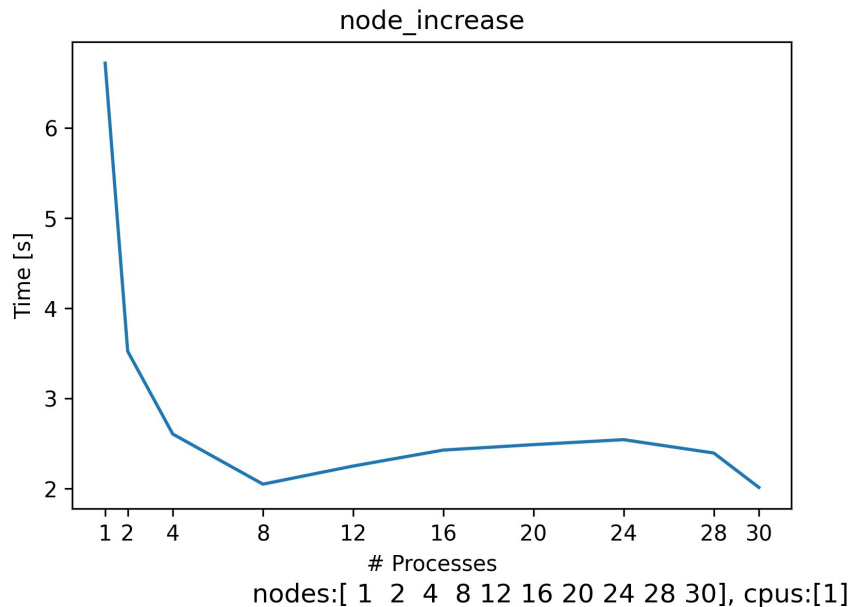
# Results Experiment 1



## Results Experiment 2




# Results Experiment 3

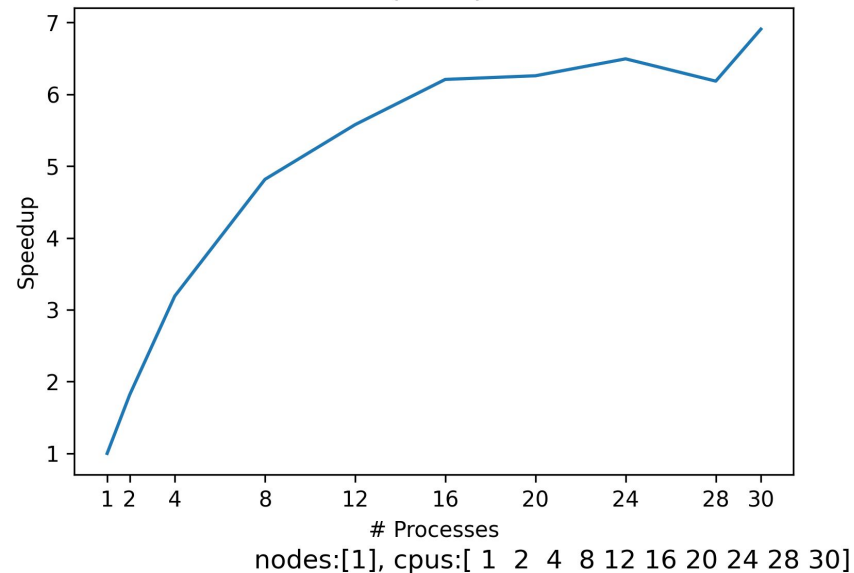


# Results Experiment 4

cpu\_increase



Speedup





# 5. Conclusion

K-Means

# Conclusion

## Learning Outcomes

- Understanding of MPI, HPC, and C

## Room for Improvement

- Improve existing parallelization
  - calculate sub-centroids in each process to then reduce to one centroid in root with custom MPI\_Op for MPI\_Reduce
- Use bigger dataset

## Further Research

- Try different parallelization approach
- Implement OpenMP version